

Skeen's protocol [1] is encoded in TLA+.

[1] *Skeen, D.*: (1985), Referenced in [1], unpublished communication.

[2] *Birman, K.P., Joseph, T.A.*: Reliable communication in the presence of failures. *ACM Transactions on Computer Systems (TOCS)* 5(1), 47–76 (1987)

Thanh-Hai *Tran, Igor Konnov, Josef Widder*, 2021 This file is a subject to the license that is bundled together with this package and can be found in the file LICENSE.

EXTENDS *Integers, FiniteSets, Sequences, TLC*

CONSTANT

@type: *Int*;
N, the number of processes indexed from 1 to *N*
 @type: *Int*;
M, the number of multicast messages indexed from 1 to *M*
 @type: *Seq(Int)*;
Mcaster, an array whose *i*-th element describes the multicaster of message *i*
 @type: *Seq(Set(Int))*;
GroupDest, an array whose *i*-th element describes the group of addressees of message *i*
 @type: *Int*;
MaxClock the bound of local clocks

All variables

VARIABLE

@type: *Int* → *Int*;
clock,
 @type: *Int* → (*Int* → *Int*);
phase,
 @type: *Int* → (*Int* → [*t*: *Int*, *g*: *Int*]);
localTS,
 @type: *Int* → (*Int* → [*t*: *Int*, *g*: *Int*]);
globalTS,
 @type: *Int* → *Set(Int)*;
rcvdMcastID,
 @type: *Set(Int)*;
mcastedID,
 @type: *Int* → (*Int* → *Set*([*type*: *Int*, *t*: *Int*, *id*: *Int*, *source*: *Int*]));
inTransit,
 @type: *Int* → (*Int* → *Bool*);
delivered,
 @type: *Int* → (*Int* → *Set*([*type*: *Int*, *t*: *Int*, *id*: *Int*, *source*: *Int*]));
proposeTS,
 @type: *Int* → (*Int* → *Int*);
dCntr

$vars \triangleq \langle clock, phase, localTS, globalTS, rcvdMcastID, mcastedID, inTransit, delivered, proposeTS, dCntr \rangle$

$Proc \triangleq 1 \dots N$
 $McastID \triangleq 1 \dots M$
 $MType \triangleq 10$ type of multicast messages
 $PType \triangleq 11$ type of proposed messages
 $Start \triangleq 12$
 $Proposed \triangleq 13$
 $Committed \triangleq 14$

ASSUME

$\wedge N \in Int$
 $\wedge \forall id \in McastID : GroupDest[id] \in SUBSET Int$
 $\wedge MType \in Int$
 $\wedge PType \in Int$
 $\wedge Start \in Int$
 $\wedge Proposed \in Int$
 $\wedge Committed \in Int$
 $\wedge M = Cardinality(McastID)$

$McastMsgPhase \triangleq \{Start, Proposed, Committed\}$

$McastPhase \triangleq [McastID \rightarrow McastMsgPhase]$

TimestampNull: the init value of local timestamps and global timestamps

Type of *TimestampNull* is $[t \mapsto Int, g \mapsto Int]$

$GroupNull \triangleq 0$

$TimeNull \triangleq 0$

type: $[t : Int, g : Int]$

$TimestampNull \triangleq [t \mapsto TimeNull, g \mapsto GroupNull]$

type: $Set(Int)$

$Time \triangleq 1 \dots MaxClock$

type: $Set(Int)$

$ProcWithNull \triangleq 0 \dots N$

The set of all possible in-transit messages

@type: $Set([t : Int, g : Int])$;

$TimestampSet \triangleq [t : Time, g : Proc] \cup \{TimestampNull\}$

@type: $Set([type : Int, t : Int, id : Int, source : Int])$;

$McastMsgSet \triangleq [t : Time, id : McastID, type : \{MType\}, source : Proc]$

@type: $Set([type : Int, t : Int, id : Int, source : Int])$;

$ProposeMsgSet \triangleq [t : Time, id : McastID, type : \{PType\}, source : Proc]$

@type: $Set([type : Int, t : Int, id : Int, source : Int])$;

$InTransitMsgSet \triangleq McastMsgSet \cup ProposeMsgSet$

The initialized states

- clock: local clocks
- $phase[p][m]$: stores the status of message m at process p
- $localTS[p][m]$: stores the local timestamp issued by process p for message m
- $globalTS[p][m]$: stores the global timestamp issued by process p for message m
- $delivered[p][m]$: refers to whether process p has delivered message m
- $rcvdMcastID[p][m]$: a set of multicast messages that process p has received
- $proposeTS[p][m]$: stores a set of proposals for messages m
- $mcastedID$: a set of messages that were multicast
- $inTransit[p][q]$: a set of in-transit messages from process p to process q
- $dCntr[p][m]$ to keep track of how many times process p has delivered message m .

$Init \triangleq$

$$\begin{aligned}
& \wedge clock = [p \in Proc \mapsto 0] \\
& \wedge phase = [p \in Proc \mapsto [m \in McastID \mapsto Start]] \\
& \wedge localTS = [p \in Proc \mapsto [m \in McastID \mapsto TimestampNull]] \\
& \wedge globalTS = [p \in Proc \mapsto [m \in McastID \mapsto TimestampNull]] \\
& \wedge delivered = [p \in Proc \mapsto [m \in McastID \mapsto FALSE]] \\
& \wedge rcvdMcastID = [p \in Proc \mapsto \{\}] \\
& \wedge proposeTS = [p \in Proc \mapsto [id \in McastID \mapsto \{\}]] \\
& \wedge mcastedID = \{\} \\
& \wedge inTransit = [p \in Proc \mapsto [q \in Proc \mapsto \{\}]] \\
& \wedge dCntr = [p \in Proc \mapsto [id \in McastID \mapsto 0]]
\end{aligned}$$

$Max(a, b) \triangleq \text{IF } a > b \text{ THEN } a \text{ ELSE } b$

Process $snder$ multicasts the message whose identifier is mid .

The multicast message for message mid is tag with a local timestamp issued by process $snder$.

$Multicast(mid) \triangleq$

LET $snder \triangleq Mcaster[mid]$

IN $\wedge mid \notin mcastedID$

$\wedge clock[snder] < MaxClock$

Only for bounded model checking

$\wedge snder \in GroupDest[mid]$

$\wedge mcastedID' = mcastedID \cup \{mid\}$

Marks that message mid is multicast

$\wedge \text{LET } time \triangleq clock[snder] + 1$

@type: [type: Int , t : Int , id : Int , source: Int];

$msg \triangleq [type \mapsto MType, id \mapsto mid, t \mapsto time, source \mapsto snder]$

IN $\wedge inTransit' = [p \in Proc \mapsto [q \in Proc \mapsto$

IF $p = snder \wedge q \in GroupDest[mid]$

THEN $inTransit[p][q] \cup \{msg\}$

ELSE $inTransit[p][q]$]

$\wedge clock' = [clock \text{ EXCEPT } ![snder] = time]$

$\wedge \text{UNCHANGED } \langle phase, proposeTS, rcvdMcastID, localTS, globalTS, delivered, dCntr \rangle$

Pick the in-transit message with the smallest timestamp from process $snder$ to process $rcver$

@type: (Int, Int , [type: Int , t : Int , id : Int , source: Int]) $\Rightarrow Bool$;

$isYoungestMsg(snder, rcver, msg) \triangleq$
 $\forall m \in inTransit[snder][rcver] : msg.t \leq m.t$

Receives a multicast message

$ReceiveMulticast(snder, rcver, msg) \triangleq$
 $\wedge clock[rcver] < MaxClock$
 $\wedge msg.type = MType$
 $\wedge isYoungestMsg(snder, rcver, msg)$ msg must have the smallest timestamp in $inTransit[snder][rcver]$
 $\wedge rcvdMcastID' = [rcvdMcastID \text{ EXCEPT } ![rcver] = rcvdMcastID[rcver] \cup \{msg.id\}]$
 $\wedge \text{UNCHANGED } \langle proposeTS, globalTS, delivered, mcastedID, dCntr \rangle$
 $\wedge \text{LET } mid \triangleq msg.id$
 $\quad time \triangleq clock[rcver] + 1$
 $\quad newTS \triangleq [t \mapsto time, g \mapsto rcver]$ the local timestamp for message $msg.id$
 $\quad @type: [type: Int, t: Int, id: Int, source: Int];$
 $\quad newMsg \triangleq [type \mapsto PType, id \mapsto mid, source \mapsto rcver, t \mapsto time]$ the proposal for message $msg.id$
IN $\wedge clock' = [clock \text{ EXCEPT } ![rcver] = clock[rcver] + 1]$
 $\wedge localTS' = [localTS \text{ EXCEPT } ![rcver][mid] = newTS]$
 $\wedge phase' = [phase \text{ EXCEPT } ![rcver][mid] = Proposed]$
Sends its proposal to every addressee of message $msg.id$
 $\wedge \text{IF } snder \neq rcver$
 $\quad \text{THEN } inTransit' = [p \in Proc \mapsto [q \in Proc \mapsto$
 $\quad \quad \text{IF } p = rcver \wedge q \in GroupDest[mid]$
 $\quad \quad \quad \text{THEN } inTransit[p][q] \cup \{newMsg\}$
 $\quad \quad \quad \text{ELSE IF } p = snder \wedge q = rcver$
 $\quad \quad \quad \quad \text{THEN } inTransit[p][q] \setminus \{msg\}$
 $\quad \quad \quad \quad \text{ELSE } inTransit[p][q]]]$
 $\quad \text{ELSE } inTransit' = [p \in Proc \mapsto [q \in Proc \mapsto$
 $\quad \quad \text{IF } p = rcver \wedge q = rcver$
 $\quad \quad \quad \text{THEN } (inTransit[p][q] \cup \{newMsg\}) \setminus \{msg\}$
 $\quad \quad \quad \text{ELSE IF } p = rcver \wedge q \in GroupDest[mid]$
 $\quad \quad \quad \quad \text{THEN } inTransit[p][q] \cup \{newMsg\}$
 $\quad \quad \quad \quad \text{ELSE } inTransit[p][q]]]$

Compare two timestamps based on lexicographical order

$@type: ([t: Int, g: Int], [t: Int, g: Int]) \Rightarrow Bool;$
 $Less(ts1, ts2) \triangleq$
 $\vee ts1.t < ts2.t$
 $\vee \wedge ts1.t = ts2.t$
 $\quad \wedge ts1.g < ts2.g$

Check whether message id can be delivered to process p

The local timestamps of all committed messages must be greater than the global timestamp of message id
 $@type: (Int, Int) \Rightarrow Bool;$

$CanDeliver(p, id) \triangleq$
 $\wedge \neg delivered[p][id]$
 $\wedge phase'[p][id] = Committed$
 $\wedge \forall mid \in rcvdMcastID'[p] :$
 $phase'[p][mid] = Proposed \Rightarrow Less(globalTS'[p][id], localTS'[p][mid])$

Process *rcver* has received the proposals from all addressees of message *id*.
 $HasAllProposes(rcver, id) \triangleq$
 $\forall p \in GroupDest[id] : \exists m \in proposeTS'[rcver][id] : m.source = p$

Pick a proposed message with the greatest local timestamp for message *id*
 $PickMsgWithMaxTS(rcver, id) \triangleq$
 CHOOSE $m \in proposeTS'[rcver][id] :$
 $\forall m1 \in proposeTS'[rcver][id] :$
 $\vee m1.t < m.t$
 $\vee \wedge m1.t = m.t$
 $\wedge m1.source \leq m.source$

Process *rcver* has received a proposed message from process *snder*
 $ReceivePropose(snder, rcver, msg) \triangleq$
 $\wedge msg.type = PType$
 $\wedge isYoungestMsg(snder, rcver, msg)$ msg must have the smallest timestamp in $inTransit[(snder, rcver)]$
 $\wedge inTransit' = [inTransit \text{ EXCEPT } ![snder][rcver] = inTransit[snder][rcver] \setminus \{msg\}]$
 $\wedge LET \ ts \triangleq [t \mapsto msg.t, g \mapsto msg.source]$
 $\quad id \triangleq msg.id$
 IN $\wedge UNCHANGED \langle localTS, mcastedID, rcvdMcastID \rangle$
 $\wedge proposeTS' = [proposeTS \text{ EXCEPT } ![rcver][id] = proposeTS[rcver][id] \cup \{msg\}]$
Whether process *rcver* has received the proposals from all addressees of message *id*.
 $\wedge IF \ HasAllProposes(rcver, id)$
 $\quad THEN LET \ m \triangleq PickMsgWithMaxTS(rcver, id)$
 $\quad \quad maxTS \triangleq [g \mapsto m.source, t \mapsto m.t]$
 $\quad IN$ Set the global timestamp for message *msg.id*
 $\quad \wedge \quad globalTS' = [globalTS \text{ EXCEPT } ![rcver][id] = maxTS]$
Synchronizes the local clocks
 $\quad \wedge \quad clock' = [clock \text{ EXCEPT } ![rcver] = Max(clock[rcver], maxTS.t)]$
 $\quad \wedge \quad phase' = [phase \text{ EXCEPT } ![rcver][id] = Committed]$
 $\quad \wedge \quad delivered' = [delivered \text{ EXCEPT } ![rcver] = [mid \in McastID \mapsto$
 $\quad \quad IF \ \neg delivered[rcver][mid]$
 $\quad \quad THEN \ CanDeliver(rcver, mid)$
 $\quad \quad ELSE \ delivered[rcver][mid]]]$
Update how many times *p* has delivered message *mid*
 $\quad \wedge \quad dCntr' = [dCntr \text{ EXCEPT } ![rcver] = [mid \in McastID \mapsto$
 $\quad \quad IF \ \neg delivered[rcver][mid] \wedge CanDeliver(rcver, mid)$
 $\quad \quad THEN \ dCntr[rcver][mid] + 1$
 $\quad \quad ELSE \ dCntr[rcver][mid]]]$
 $\quad ELSE \ UNCHANGED \langle phase, globalTS, clock, delivered, dCntr \rangle$

Only to avoid deadlock checking

$Done \triangleq$

$\wedge \forall id \in McastID : \forall p \in GroupDest[id] : delivered[p][id]$
 $\wedge UNCHANGED\ vars$

$Next \triangleq$

$\vee \exists m \in McastID : Multicast(m)$
 $\vee \exists sndr, rcvr \in Proc : \exists msg \in inTransit[sndr][rcvr] : ReceiveMulticast(sndr, rcvr, msg)$
 $\vee \exists sndr, rcvr \in Proc : \exists msg \in inTransit[sndr][rcvr] : ReceivePropose(sndr, rcvr, msg)$
 $\vee Done$

$Spec \triangleq$

$\wedge Init$
 $\wedge \Box [Next]_{vars}$
 $\wedge WF_{vars}(\vee \exists m \in McastID : Multicast(m)$
 $\vee \exists sndr, rcvr \in Proc : \exists msg \in inTransit[sndr][rcvr] :$
 $ReceiveMulticast(sndr, rcvr, msg)$
 $\vee \exists sndr, rcvr \in Proc : \exists msg \in inTransit[sndr][rcvr] :$
 $ReceivePropose(sndr, rcvr, msg))$

- Total Order: There exists a total order $<$ on all messages that are multicast in an execution trace such that, if process p delivers message m , then for all messages $m' < m$ such that p is one of addresses of message m' , p delivers m' before m .
- Total Order can be formalized as the following formula

$GlobalTotalOrdering \triangleq$

$\exists ordering \in [McastID \rightarrow 1..M] :$
 $\wedge \forall p \in Proc : \forall id1, id2 \in McastID :$
 $(\wedge globalTS[p][id1] \neq TimestampNull$
 $\wedge globalTS[p][id2] \neq TimestampNull$
 $\wedge ordering[id1] < ordering[id2])$
 $\Rightarrow Less(globalTS[p][id1], globalTS[p][id2])$

- However, APALACHE cannot verify $GlobalTotalOrdering$ because the initialization of ordering and its corresponding quantifiers.

The conjunction of $ConsistentGlobalTS$ and $AsymmetricOrdering$ implies Total Order

$AsymmetricOrdering \triangleq$

$\forall id1, id2 \in McastID : \forall p, q \in Proc :$
 $(\wedge globalTS[p][id1] \neq TimestampNull$
 $\wedge globalTS[p][id2] \neq TimestampNull$
 $\wedge globalTS[q][id1] \neq TimestampNull$
 $\wedge globalTS[q][id2] \neq TimestampNull$
 $\wedge id1 \neq id2)$
 $\Rightarrow \neg(Less(globalTS[p][id1], globalTS[p][id2]) \wedge Less(globalTS[q][id2], globalTS[q][id1]))$

$ConsistentGlobalTS \triangleq$

$\wedge \forall id \in McastID : \forall p, q \in Proc :$

All addressees of message id must agree on its global timestamp.

$$\begin{aligned}
& (\wedge globalTS[p][id] \neq TimestampNull \\
& \quad \wedge globalTS[q][id] \neq TimestampNull) \\
& \quad \Rightarrow globalTS[p][id] = globalTS[q][id] \\
\wedge \forall id1, id2 \in McastID : \forall p \in Proc : & \quad \text{Every message has a unique global timestamp.} \\
& (\wedge id1 \neq id2 \\
& \quad \wedge globalTS[p][id1] \neq TimestampNull \\
& \quad \wedge globalTS[p][id2] \neq TimestampNull) \\
& \quad \Rightarrow globalTS[p][id1] \neq globalTS[p][id2]
\end{aligned}$$

$$Validity \triangleq \forall p \in Proc : \forall id \in McastID : delivered[p][id] \Rightarrow id \in mcastedID$$

If process p is not an addressee of message id , p never issues a local timestamp for id .
 Process p issues a local timestamp for message id if and only if it receive a multicast message for id .
 The time in every local timestamp cannot greater than the current value of the local clock.
 Never issues a local timestamp with *GroupNull* or *TimestampNull*.
 The owner of the local timestamp $localTS[p][id]$ must be process p .
 Never issues two local timestamps at one time point.

$$\begin{aligned}
ValidOwnedLocalTS & \triangleq \\
& \wedge (\forall id \in McastID : \forall p \in Proc \setminus GroupDest[id] : localTS[p][id] = TimestampNull) \\
& \wedge (\forall id \in McastID : \forall p \in GroupDest[id] : \\
& \quad \wedge localTS[p][id] = TimestampNull \equiv id \notin rcvdMcastID[p] \\
& \quad \wedge localTS[p][id].t \leq clock[p] \\
& \quad \wedge (localTS[p][id].g \neq GroupNull \Rightarrow localTS[p][id].t \neq TimeNull) \\
& \quad \wedge (localTS[p][id] \neq TimestampNull \\
& \quad \quad \Rightarrow (\wedge id \in mcastedID \\
& \quad \quad \quad \wedge localTS[p][id].g = p))) \\
& \wedge \forall id1, id2 \in McastID : \forall p \in Proc : \\
& \quad ((\wedge p \in GroupDest[id1] \\
& \quad \quad \wedge p \in GroupDest[id2] \\
& \quad \quad \wedge id1 \neq id2 \\
& \quad \quad \wedge localTS[p][id1] \neq TimestampNull \\
& \quad \quad \wedge localTS[p][id2] \neq TimestampNull) \\
& \quad \quad \Rightarrow localTS[p][id1].t \neq localTS[p][id2].t)
\end{aligned}$$

Every in-transit message in $inTransit[snder][rcver]$ was sent by process $snder$.
 The in-transit proposed message for message id must be sent after the multicast message for message id .

$$\begin{aligned}
ValidInTransitMsg & \triangleq \\
& \wedge \forall snder, rcver \in Proc : \forall m \in inTransit[snder][rcver] : m.source = snder \\
& \wedge \forall snder, rcver \in Proc : \forall m1, m2 \in inTransit[snder][rcver] : \\
& \quad ((\wedge m1.id = m2.id \\
& \quad \quad \wedge m1.type = MType \\
& \quad \quad \wedge m2.type = PType)
\end{aligned}$$

$$\Rightarrow m1.t < m2.t)$$

- If process p is not an addressee of message id , no processes send a proposal for message id to process p .
- If process p is not an addressee of message id , it never sends a proposal for message id .
- If there exists a proposal for message id from process $snder$, process $snder$ has issued a local timestamp for message m . These timestamps must be the same.
- If there exists a proposal for message id , message id must be multicast before.
- The time in an issued timestamp by process $snder$ cannot greater than the current value of the clock of process $snder$.
- If there exists an in-transit proposed message for message id that is sent to process $rcver$, process $rcver$ has not issued a global timestamp for message id .
- If there exists an in-transit proposed message for message id that is sent from process $snder$, process $snder$ has issued a local timestamp for message id .
- If there exists an in-transit proposed message for message id , message id must be multicast before.
- If there exists an in-transit proposed message for message id that is sent from process $snder$, there exists no in-transit multicast message to process $snder$ such that this multicast message is for message id .

$ValidInTransitProposeTS \triangleq$

$$\begin{aligned}
& \wedge (\forall id \in McastID : \forall rcver \in Proc \setminus GroupDest[id] : \forall snder \in Proc : \\
& \quad \forall m \in inTransit[snder][rcver] : m.id \neq id) \\
& \wedge (\forall id \in McastID : \forall snder \in Proc \setminus GroupDest[id] : \forall rcver \in Proc : \\
& \quad \forall m \in inTransit[snder][rcver] : m.id \neq id) \\
& \wedge (\forall id \in McastID : \forall rcver \in GroupDest[id] : \forall snder \in Proc : \forall m \in inTransit[snder][rcver] : \\
& \quad (\wedge m.id = id \\
& \quad \wedge m.source = snder \\
& \quad \wedge m.type = PType) \\
& \quad \Rightarrow (\wedge m.t = localTS[snder][id].t \\
& \quad \wedge id \in rcvdMcastID[snder] \\
& \quad \wedge id \in mcastedID)) \\
& \wedge (\forall snder, rcver \in Proc : \forall m \in inTransit[snder][rcver] : m.t \leq clock[m.source] \wedge m.t > TimeNull) \\
& \wedge (\forall snder, rcver \in Proc : \forall m \in inTransit[snder][rcver] : \\
& \quad m.t = PType \Rightarrow (\wedge globalTS[rcver][m.id] = TimestampNull \\
& \quad \wedge localTS[m.source][m.id] \neq TimestampNull \\
& \quad \wedge m.id \in rcvdMcastID[m.source])) \\
& \wedge (\forall snder, rcver \in Proc : \forall m \in inTransit[snder][rcver] : \\
& \quad m.t = PType \Rightarrow (\wedge localTS[m.source][m.id].g = m.source \\
& \quad \wedge localTS[m.source][m.id].t = m.t)) \\
& \wedge (\forall id \in McastID : \forall snder, rcver \in GroupDest[id] : \forall m \in inTransit[snder][rcver] : \\
& \quad ((m.t = PType \wedge m.id = id) \\
& \quad \Rightarrow (\forall m1 \in inTransit[Mcaster[id]][snder] : m1.type = MType \Rightarrow m1.id \neq id)))
\end{aligned}$$

- If process p is not an addressee of message id , it never receives a proposal for message id .
- Every received proposed message for message id must be grouped correctly.
- Every received proposed message for message id from process $snder$ must propose the local timestamp that is issued by process $snder$ for message id .

$$\begin{aligned}
\text{ValidRcvdProposeTS} &\triangleq \\
&\wedge (\forall id \in \text{McastID} : \forall rcver \in \text{Proc} \setminus \text{GroupDest}[id] : \forall sndr \in \text{Proc} : \\
&\quad \text{proposeTS}[rcver][id] = \{\}) \\
&\wedge (\forall id \in \text{McastID} : \forall rcver \in \text{GroupDest}[id] : \forall msg \in \text{proposeTS}[rcver][id] : \\
&\quad \wedge msg.t = \text{localTS}[msg.source][msg.id].t \\
&\quad \wedge msg.id = id \\
&\quad \wedge (\forall m \in \text{inTransit}[msg.source][rcver] : m.id \neq id)) \\
&\wedge (\forall id \in \text{McastID} : \forall rcver \in \text{GroupDest}[id] : \forall msg \in \text{proposeTS}[rcver][id] : \\
&\quad \wedge msg.t = \text{localTS}[msg.source][msg.id].t \\
&\quad \wedge msg.source = \text{localTS}[msg.source][msg.id].g \\
&\quad \wedge msg.id = id)
\end{aligned}$$

Every local clock is bounded with *MaxClock*

$$\text{BoundedClock} \triangleq \forall p \in \text{Proc} : \text{clock}[p] \leq \text{MaxClock}$$

- The global timestamp for message m cannot be less than any proposed local timestamp for message m .
- The global timestamp for message m must equal some local timestamp for message m .
- If process p is not an addressee of message id , it never issues a global timestamp for message id .
- There exists no global timestamp with *GroupNull* or *TimeNull*.
- The time in every global timestamp cannot greater than the current value of the clock.
- The global timestamp for message id is issued if and only if message id is committed.

$$\begin{aligned}
\text{ValidGlobalTS} &\triangleq \\
&\wedge \forall id \in \text{McastID} : \forall rcver \in \text{GroupDest}[id] : \\
&\quad (\text{globalTS}[rcver][id] \neq \text{TimestampNull} \\
&\quad \equiv (\wedge \forall sndr \in \text{GroupDest}[id] : \exists m \in \text{proposeTS}[rcver][id] : \\
&\quad \quad (\wedge m.source = sndr \\
&\quad \quad \wedge \forall m.t < \text{globalTS}[rcver][id].t \\
&\quad \quad \vee \wedge m.t = \text{globalTS}[rcver][id].t \\
&\quad \quad \wedge m.source \leq \text{globalTS}[rcver][id].g)) \\
&\quad \wedge \exists sndr \in \text{GroupDest}[id] : \exists m \in \text{proposeTS}[rcver][id] : \\
&\quad \quad (\wedge \text{globalTS}[rcver][id].t = m.t \\
&\quad \quad \wedge \text{globalTS}[rcver][id].g = m.source)) \\
&\wedge \forall id \in \text{McastID} : \forall rcver \in \text{Proc} \setminus \text{GroupDest}[id] : \text{globalTS}[rcver][id] = \text{TimestampNull} \\
&\wedge \forall id \in \text{McastID} : \forall rcver \in \text{GroupDest}[id] : \\
&\quad \text{globalTS}[rcver][id].g \neq \text{GroupNull} \Rightarrow \text{globalTS}[rcver][id].t \neq \text{TimeNull} \\
&\wedge \forall id \in \text{McastID} : \forall p \in \text{Proc} \setminus \text{GroupDest}[id] : \text{globalTS}[p][id] = \text{TimestampNull} \\
&\wedge \forall id \in \text{McastID} : \forall p \in \text{GroupDest}[id] : \\
&\quad \text{globalTS}[p][id] \neq \text{TimestampNull} \Rightarrow \exists q \in \text{GroupDest}[id] : \text{localTS}[q][id] = \text{globalTS}[p][id] \\
&\wedge \forall id \in \text{McastID} : \forall rcver \in \text{GroupDest}[id] : \\
&\quad \text{globalTS}[rcver][id].g \neq \text{GroupNull} \Rightarrow \text{globalTS}[rcver][id].t \leq \text{clock}[rcver] \\
&\wedge \forall id \in \text{McastID} : \forall p \in \text{Proc} : \text{globalTS}[p][id] \neq \text{TimestampNull} \equiv \text{phase}[p][id] = \text{Committed}
\end{aligned}$$

Process p sets the status of message id to *Start* iff it has not issued a local timestamp for message id .

If process p commits message id , it has received at least one proposal for message id .

If process p commits message id , it has not issued any global timestamp for message id .

$ValidPhase \triangleq$

$$\begin{aligned} & \forall p \in Proc : \forall id \in McastID : \\ & \quad (\wedge p \notin GroupDest[id] \Rightarrow phase[p][id] = Start \\ & \quad \wedge phase[p][id] = Start \equiv localTS[p][id] = TimestampNull \\ & \quad \wedge phase[p][id] = Proposed \Rightarrow (localTS[p][id] \neq TimestampNull \wedge id \in rcvdMcastID[p]) \\ & \quad \wedge phase[p][id] = Committed \equiv (\forall q \in GroupDest[id] : \exists m \in proposeTS[p][id] : m.source = q) \\ & \quad \wedge ((localTS[p][id] \neq TimestampNull \wedge id \in rcvdMcastID[p]) \\ & \quad \Rightarrow phase[p][id] \in \{Proposed, Committed\})) \end{aligned}$$

Message id can be delivered to process p if and only if process p has issued a global timestamp for message id

and the local timestamps of all proposed message at process p must be greater than the global timestamp of message id .

$ValidDelivery \triangleq$

$$\begin{aligned} & \forall p \in Proc : \forall id \in McastID : \\ & \quad delivered[p][id] \\ & \Rightarrow (\wedge globalTS[p][id] \neq TimestampNull \\ & \quad \wedge phase[p][id] = Committed \\ & \quad \wedge \forall mid \in rcvdMcastID[p] : \\ & \quad \quad phase[p][mid] = Proposed \Rightarrow Less(globalTS[p][id], localTS[p][mid])) \end{aligned}$$

Every in-transit message has an unique timestamp.

If process $snder$ has sent a proposal for message id , no in-transit message to process p is a multicast message for message id .

$UniqueMsg \triangleq$

$$\begin{aligned} & \wedge (\forall snder, rcver \in Proc : \forall m1, m2 \in inTransit[snder][rcver] : \\ & \quad \wedge (m1.type = m2.type \wedge m1.id = m2.id) \Rightarrow m1.t = m2.t \\ & \quad \wedge (m1.type = m2.type \wedge m1.t = m2.t) \Rightarrow m1.id = m2.id \\ & \quad \wedge (m1.id = m2.id \wedge m1.t = m2.t) \Rightarrow m1.type = m2.type) \\ & \wedge (\forall snder, rcver \in Proc : \forall m \in inTransit[snder][rcver] : \\ & \quad m.type = PType \Rightarrow \neg(\exists m1 \in inTransit[Mcaster[m.id]][snder] : \\ & \quad \quad \wedge m1.id = m.id \\ & \quad \quad \wedge m1.source = Mcaster[m.id] \\ & \quad \quad \wedge m1.type = MType)) \end{aligned}$$

- If process p is not an addressee of message id , it never receives a multicast message for message id .

- Every multicast message for message id must be multicast by its multicaster.

- If there exists a multicast message for message id , message id must be multicast before.

- The timestamp of every proposed message from process $snder$ cannot be greater the local clock of process $snder$, and must be not 0.

- If there exists an in-transit multicast message for message id to process $rcver$, process $rcver$ has not issued neither local timestamp nor global timestamp for message id .

$ValidInTransitMcast \triangleq$

$$\wedge \forall snder, rcver \in Proc : \forall id \in McastID : \forall m \in inTransit[snder][rcver] :$$

$$\begin{aligned}
& (m.type = MType \wedge m.id = id) \Rightarrow (sndr = Mcaster[id] \wedge id \in mcastedID) \\
& \wedge \forall sndr, rcver \in Proc : \forall m \in inTransit[sndr][rcver] : \\
& \quad m.type = MType \Rightarrow m.source = Mcaster[m.id] \\
& \wedge \forall sndr, rcver \in Proc : \forall m \in inTransit[sndr][rcver] : \\
& \quad m.type = MType \Rightarrow m.id \in mcastedID \\
& \wedge \forall sndr, rcver \in Proc : \forall m \in inTransit[sndr][rcver] : \\
& \quad m.t \leq clock[m.source] \wedge m.t > TimeNull \\
& \wedge \forall mcaster, rcver \in Proc : \forall m \in inTransit[mcaster][rcver] : \\
& \quad m.type = MType \Rightarrow (\wedge \neg(\exists q \in Proc : \exists m1 \in inTransit[rcver][q] : \\
& \quad \quad \wedge m1.source = rcver \\
& \quad \quad \wedge m1.id = m.id \\
& \quad \quad \wedge m1.type = PType) \\
& \quad \wedge localTS[rcver][m.id] = TimestampNull \\
& \quad \wedge \forall p \in GroupDest[m.id] : globalTS[p][m.id] = TimestampNull)
\end{aligned}$$

$$\begin{aligned}
TypeOK & \triangleq \\
& \wedge clock \in [Proc \rightarrow Time \cup \{TimeNull\}] \\
& \wedge localTS \in [Proc \rightarrow [McastID \rightarrow TimestampSet]] \\
& \wedge globalTS \in [Proc \rightarrow [McastID \rightarrow TimestampSet]] \\
& \wedge phase \in [Proc \rightarrow [McastID \rightarrow \{Start, Proposed, Committed\}]] \\
& \wedge rcvdMcastID \in [Proc \rightarrow SUBSET McastID] \\
& \wedge mcastedID \in SUBSET McastID \\
& \wedge inTransit \in [Proc \rightarrow [Proc \rightarrow SUBSET InTransitMsgSet]] \\
& \wedge delivered \in [Proc \rightarrow [McastID \rightarrow BOOLEAN]] \\
& \wedge proposeTS \in [Proc \rightarrow [McastID \rightarrow SUBSET ProposeMsgSet]] \\
& \wedge dCntr \in [Proc \rightarrow [McastID \rightarrow \{0, 1\}]]
\end{aligned}$$

$$\begin{aligned}
Integrity & \triangleq \\
& \forall id \in McastID : \forall p \in Proc : \\
& \quad \wedge delivered[p][id] \equiv dCntr[p][id] = 1 \\
& \quad \wedge \neg delivered[p][id] \equiv dCntr[p][id] = 0
\end{aligned}$$

$$\begin{aligned}
IndInv & \triangleq \\
& \wedge TypeOK \\
& \wedge Integrity \\
& \wedge Validity \\
& \wedge ValidInTransitMsg \\
& \wedge AsymmetricOrdering \\
& \wedge ConsistentGlobalTS \\
& \wedge ValidOwnedLocalTS \\
& \wedge ValidInTransitProposeTS \\
& \wedge ValidRcvdProposeTS \\
& \wedge BoundedClock
\end{aligned}$$

\wedge *ValidGlobalTS*
 \wedge *ValidDelivery*
 \wedge *ValidPhase*
 \wedge *ValidInTransitMcast*
 \wedge *UniqueMsg*

* Modification History
* Last modified *Wed Oct 06 11:57:13 CEST 2021* by *tran*
* Created *Tue Mar 16 08:59:43 CET 2021* by *tran*